

# PostgreSQL Deep Dive

**SQL Basics → Indexing → Transactions → Query Planning → JSON & Extensions → Replication & Tuning**

Designed as **PDF-ready learning material** with explanations + examples.

---

## 1. SQL BASICS (FOUNDATION)

### PostgreSQL vs SQL

- **PostgreSQL**: Database engine (stores data, manages concurrency)
- **SQL**: Language used to talk to PostgreSQL

#### SQL Categories

- **DDL** (Data Definition Language): structure
  - **DML** (Data Manipulation Language): data
  - **DCL** (Data Control Language): permissions
  - **TCL** (Transaction Control Language): transactions
- 

### 1.1 Core DDL

```
CREATE TABLE accounts (
    id BIGSERIAL PRIMARY KEY,
    email TEXT UNIQUE NOT NULL,
    balance NUMERIC(12,2) DEFAULT 0,
    created_at TIMESTAMPTZ DEFAULT now()
);
```

**Why:** - `BIGSERIAL` = sequence-backed auto increment - `NUMERIC` = money-safe (no float errors) - `TIMESTAMPTZ` = timezone-safe timestamps

---

### 1.2 Core DML

```
INSERT INTO accounts(email, balance)
VALUES ('a@x.com', 1000);
```

```
UPDATE accounts SET balance = balance - 200 WHERE id = 1;  
DELETE FROM accounts WHERE id = 99;
```

PostgreSQL guarantees **ACID** for each statement.

---

## 1.3 SELECT Order of Execution

```
SELECT email, balance  
FROM accounts  
WHERE balance > 100  
ORDER BY balance DESC  
LIMIT 5;
```

Logical order: 1. FROM 2. WHERE 3. SELECT 4. ORDER BY 5. LIMIT

---

# 2. INDEXING (PERFORMANCE CORE)

## 2.1 What Indexes Really Are

Indexes = **sorted lookup structures**, not magic. Tradeoff: - Faster reads - Slower writes - More disk

---

## 2.2 B-Tree Index (Default)

```
CREATE INDEX idx_accounts_email ON accounts(email);
```

Used for: -  = -  < > -  ORDER BY

---

## 2.3 Composite Index

```
CREATE INDEX idx_accounts_email_balance  
ON accounts(email, balance);
```

Works for: - email - email + balance  NOT balance alone

---

## 2.4 Partial Index

```
CREATE INDEX idx_positive_balance  
ON accounts(balance)  
WHERE balance > 0;
```

Smaller, faster, smarter.

---

## 2.5 GIN Index (JSON / Arrays)

```
CREATE INDEX idx_data_gin  
ON products USING gin(data);
```

Required for JSONB performance.

---

# 3. TRANSACTIONS (DATA SAFETY)

## 3.1 ACID Explained

- **Atomic:** all or nothing
  - **Consistent:** rules enforced
  - **Isolated:** no dirty reads
  - **Durable:** WAL-backed
- 

## 3.2 Transaction Example

```
BEGIN;  
UPDATE accounts SET balance = balance - 500 WHERE id = 1;  
UPDATE accounts SET balance = balance + 500 WHERE id = 2;  
COMMIT;
```

If any statement fails → ROLLBACK.

---

### 3.3 Isolation Levels

```
SET TRANSACTION ISOLATION LEVEL SERIALIZABLE;
```

Levels: - READ COMMITTED (default) - REPEATABLE READ - SERIALIZABLE

Higher = safer, slower.

---

## 4. QUERY PLANNING (HOW POSTGRES THINKS)

### 4.1 EXPLAIN ANALYZE

```
EXPLAIN ANALYZE  
SELECT * FROM accounts WHERE email = 'a@x.com';
```

Look for: - Seq Scan ✗ - Index Scan ✓ - Rows vs actual rows

---

### 4.2 Common Planner Mistakes

- Missing index
- Wrong column order
- Outdated statistics

Fix with:

```
ANALYZE accounts;
```

---

### 4.3 Index-Only Scans

```
CREATE INDEX idx_covering  
ON accounts(email) INCLUDE (balance);
```

---

Avoids table reads entirely.

---

## 5. JSON & EXTENSIONS (POSTGRES SUPERPOWER)

### 5.1 JSON vs JSONB

- `json` : stored as text
- `jsonb` : binary, indexable

Always prefer **jsonb**.

---

### 5.2 JSON Queries

```
SELECT data->>'name'  
FROM products  
WHERE data @> '{"active":true}';
```

Operators: - `->` json - `->>` text - `@>` contains

---

### 5.3 Powerful Extensions

```
CREATE EXTENSION pg_trgm;    -- fuzzy search  
CREATE EXTENSION postgis;   -- geospatial  
CREATE EXTENSION citext;    -- case-insensitive text
```

Postgres = extensible database, not just SQL.

---

## 6. REPLICATION & TUNING (ADVANCED)

### 6.1 WAL (Write Ahead Log)

- Every change written to WAL first
- Guarantees durability
- Powers replication & recovery

---

### 6.2 Streaming Replication

Primary → Standby - Physical replication - Hot standby supported

Use cases: - High availability - Read scaling

---

## 6.3 Logical Replication

```
CREATE PUBLICATION pub_all FOR ALL TABLES;
CREATE SUBSCRIPTION sub
CONNECTION '...'
PUBLICATION pub_all;
```

Used for: - Version upgrades - Selective replication

---

## 6.4 Performance Tuning Essentials

Key settings: - shared\_buffers - work\_mem - maintenance\_work\_mem - effective\_cache\_size

Golden rule:

Measure first, tune second.

---

# FINAL MENTAL MODEL

- SQL = language
  - PostgreSQL = operating system for data
  - Indexes = maps
  - WAL = black box recorder
  - Planner = cost-based optimizer
- 

## You are now at Senior PostgreSQL Level

Next upgrades: - Query rewriting - Lock analysis - Vacuum internals - Sharding strategies

---

**Source:** PostgreSQL 18.x Official Documentation